

# Choosing a Deployment Model

---



**Eric Wright**

Technology Evangelist, Podcaster

@DiscoPosse [www.discoposse.com](http://www.discoposse.com) [www.discopossepodcast.com](http://www.discopossepodcast.com)

## Overview



- What makes up a Nomad environment?
- Compare deployment examples
- Explore the Wired Brain Coffee use-case
- Map out a hybrid Nomad deployment
- See Nomad example using Infrastructure-as-Code with Terraform



## **Starter Course**

### **Getting Started with HashiCorp Nomad**

Eric Wright

# Nomad Components and Dependencies

---

# Nomad Dependencies

**Ideal to have Consul as an auto-discovery + service mesh**

**Task drivers to support various jobs/workloads**

**Persistent storage plugins needed for stateful workloads**

**Secret management (e.g. Vault)**

**Monitoring for Day 2 operations**



**Compute Resources**




**Networking Resources**



**Storage Resources**


# Recommended Resource Node Minimums



**4-8+ cores**



**16-32+ GB of RAM**



**40-80 GB of “fast” disk**

# Persistent Storage for Nomad Workloads



## **Nomad Host Volumes**

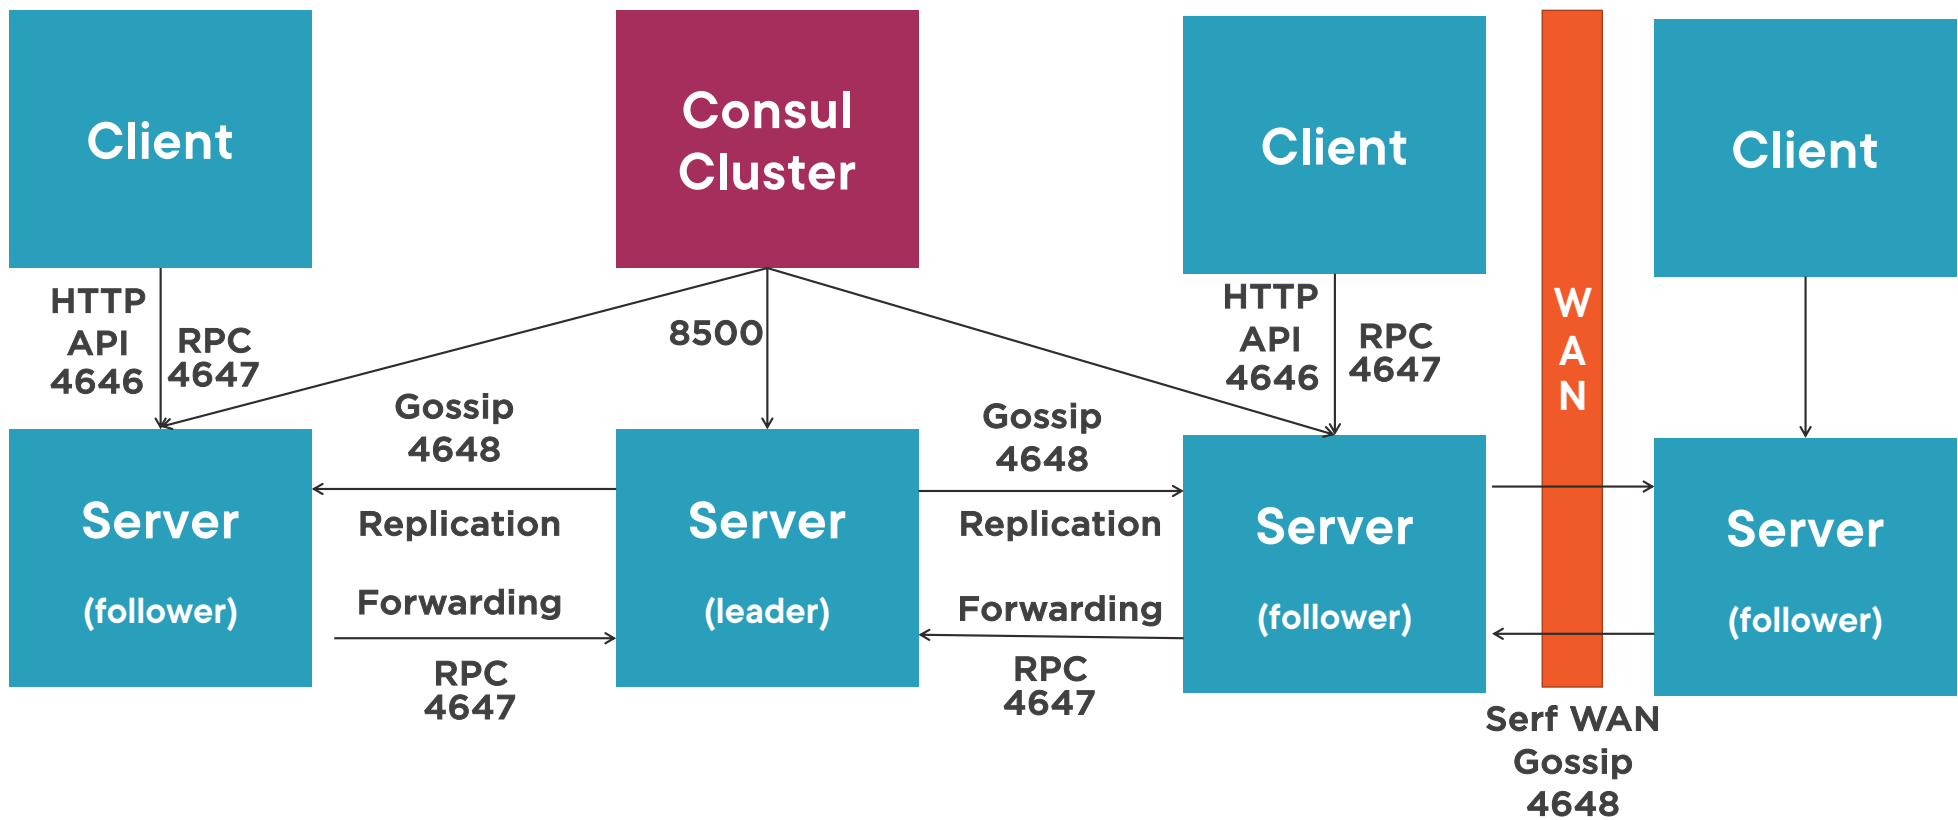
**Locally presented  
persistent storage**



## **CSI Plugins**

**Uses CSI-certified  
backend storage**

# Networking Dependencies






# Nomad Deployment Patterns

---


# Nomad Server Architectural Patterns



**Nomad expects 3 or 5 servers in a cluster**



**Nomad servers expect <10ms latency**



**Nomad clients can be deployed across any environment**

# Nomad Server and Client Permissions



Server

Nomad servers should be run with the lowest possible permissions

Create a nomad user with the minimal required privileges



Client

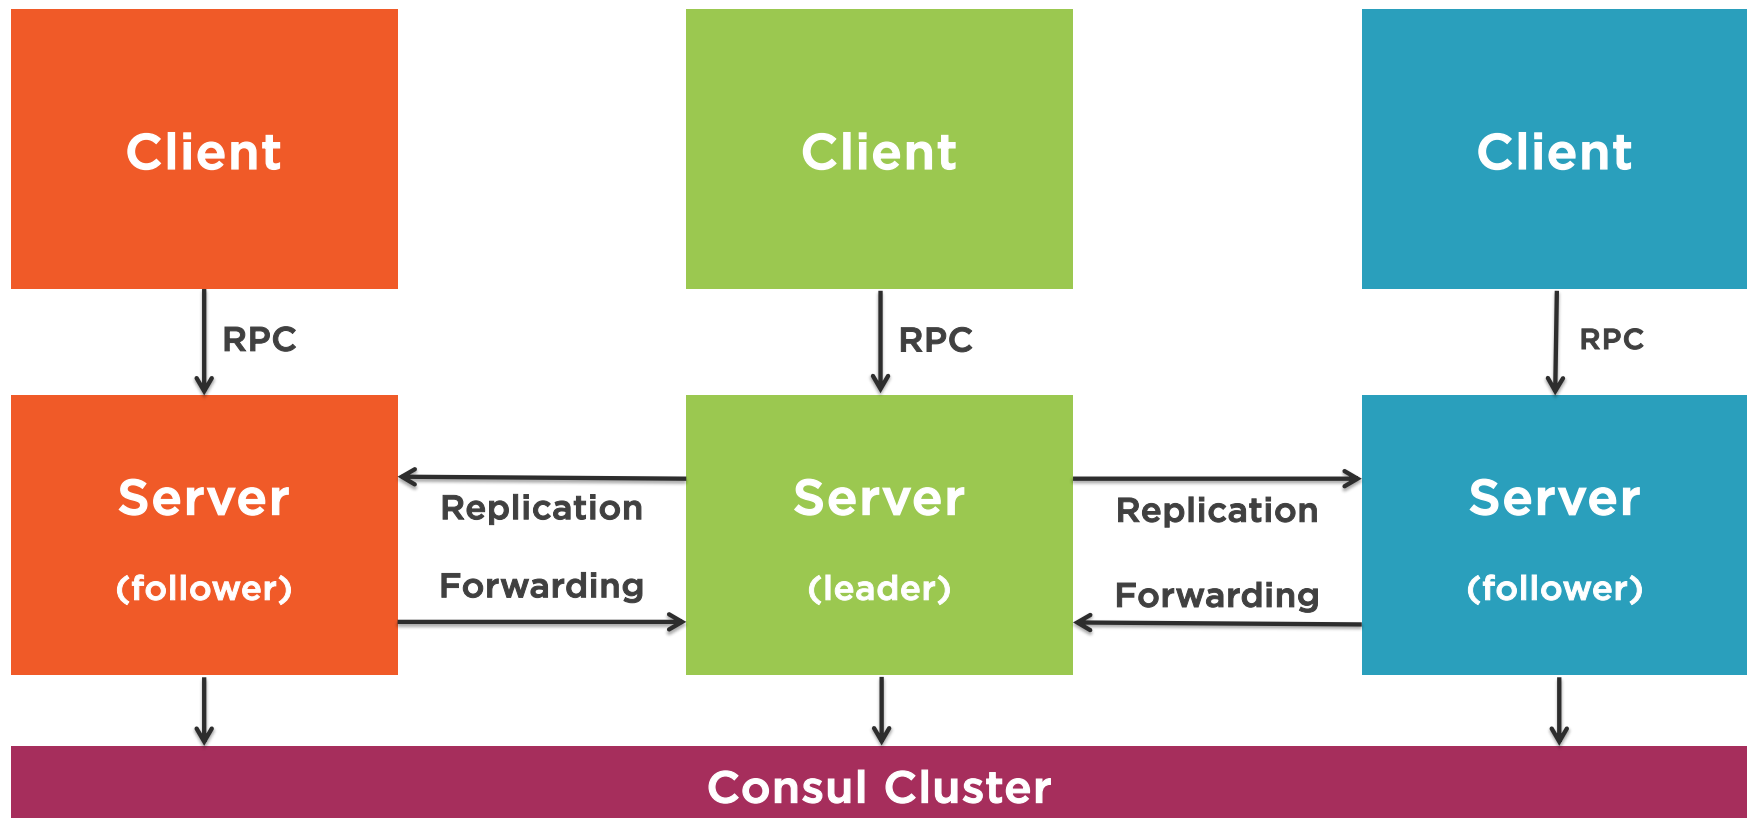
Nomad clients must be run as root due to the OS isolation

Nomad client folder needs to be owned by root (0700 permission)

# Single Region Deployment

---

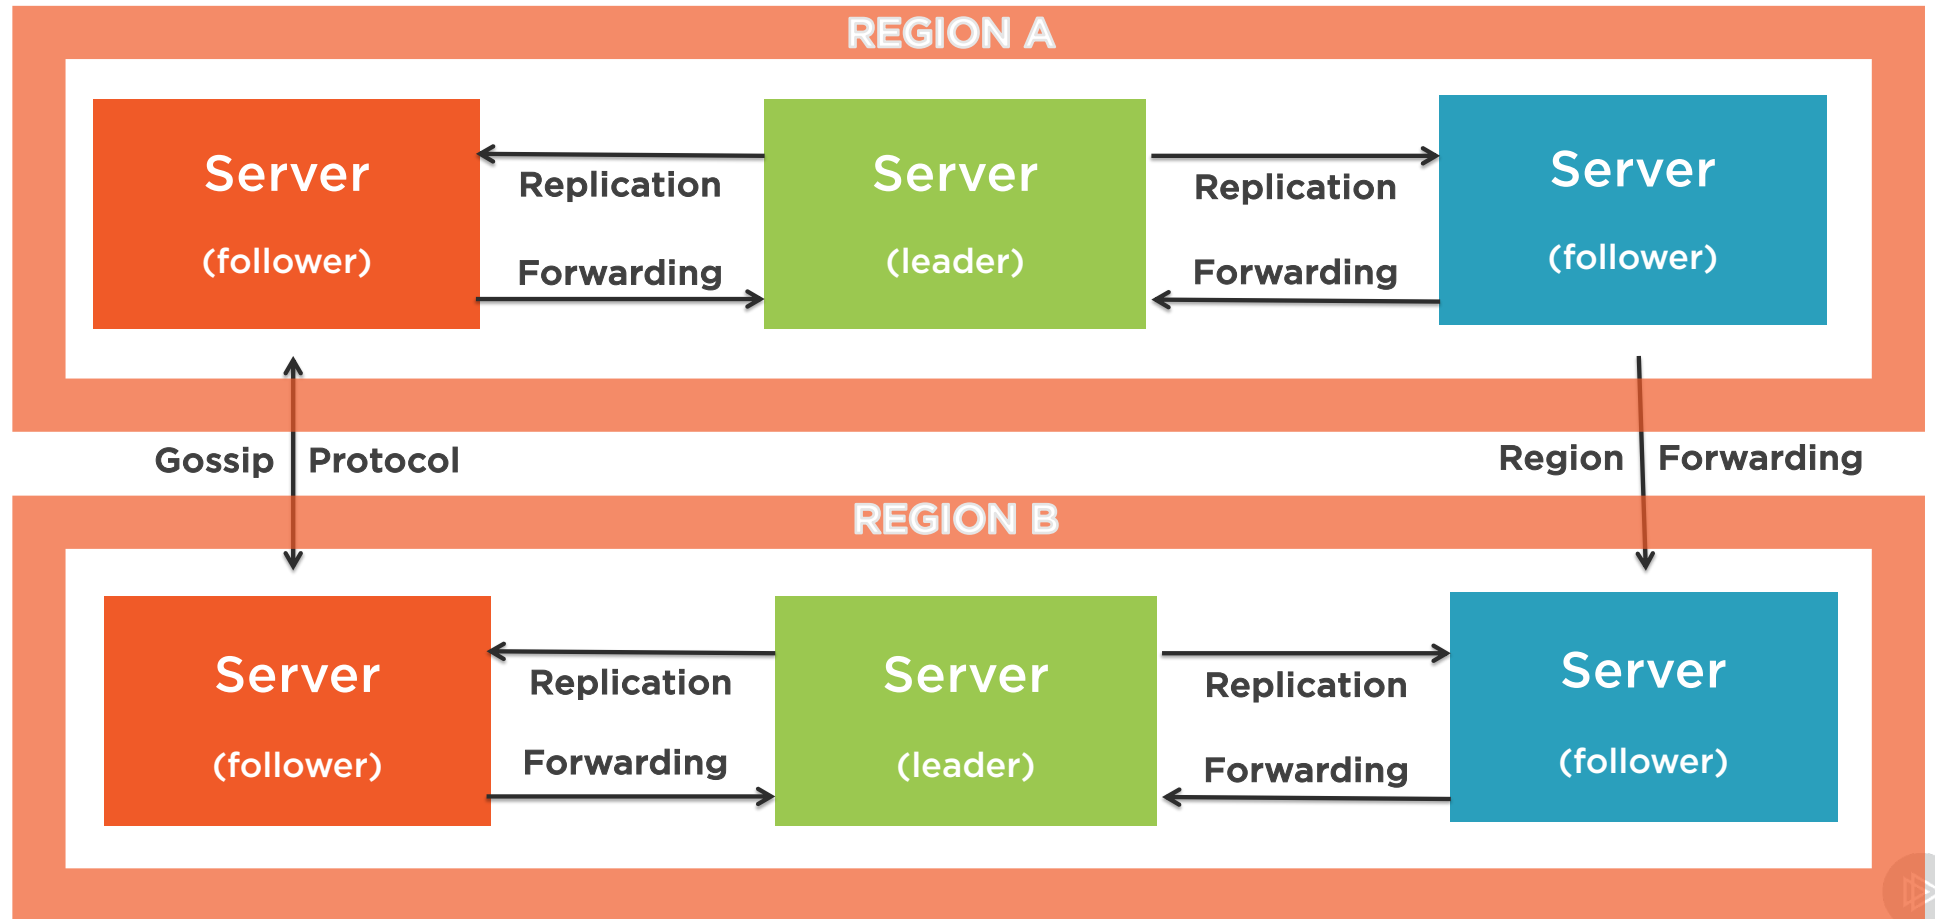
# Client/Server Deployment



# Multi-Region Deployment

---

# Multi-Region Deployment



# Understanding Nomad Namespaces

---



# Namespaces

**Nomad also supports namespaces, which allows jobs and any of their associated objects to be segmented from each other and from other users of the cluster.**

**Nomad places all job-related objects including jobs, allocations, deployments, and evaluations, using namespaces.**

# Common Namespace Use-Cases



**Segment between  
development teams**



**Segment  
environments within a  
cluster**



**Ensure co-location of  
resources for teams**


# Sizing your Nomad Deployment

---

# What Resources Need Sizing?




**# of Nodes**  
Autoscaling  
options available



**Compute**  
Requires  
common  
processors




**Memory**  
Use for rapid  
local caching



**Storage**  
Ephemeral and  
persistent

## Use-Case: Wired Brain Coffee

---



**WIRED BRAIN**  
— COFFEE —

## Application Re-Platform Project

You are the Ops team lead at Wired Brain Coffee you have chosen Nomad as your hosting platform.

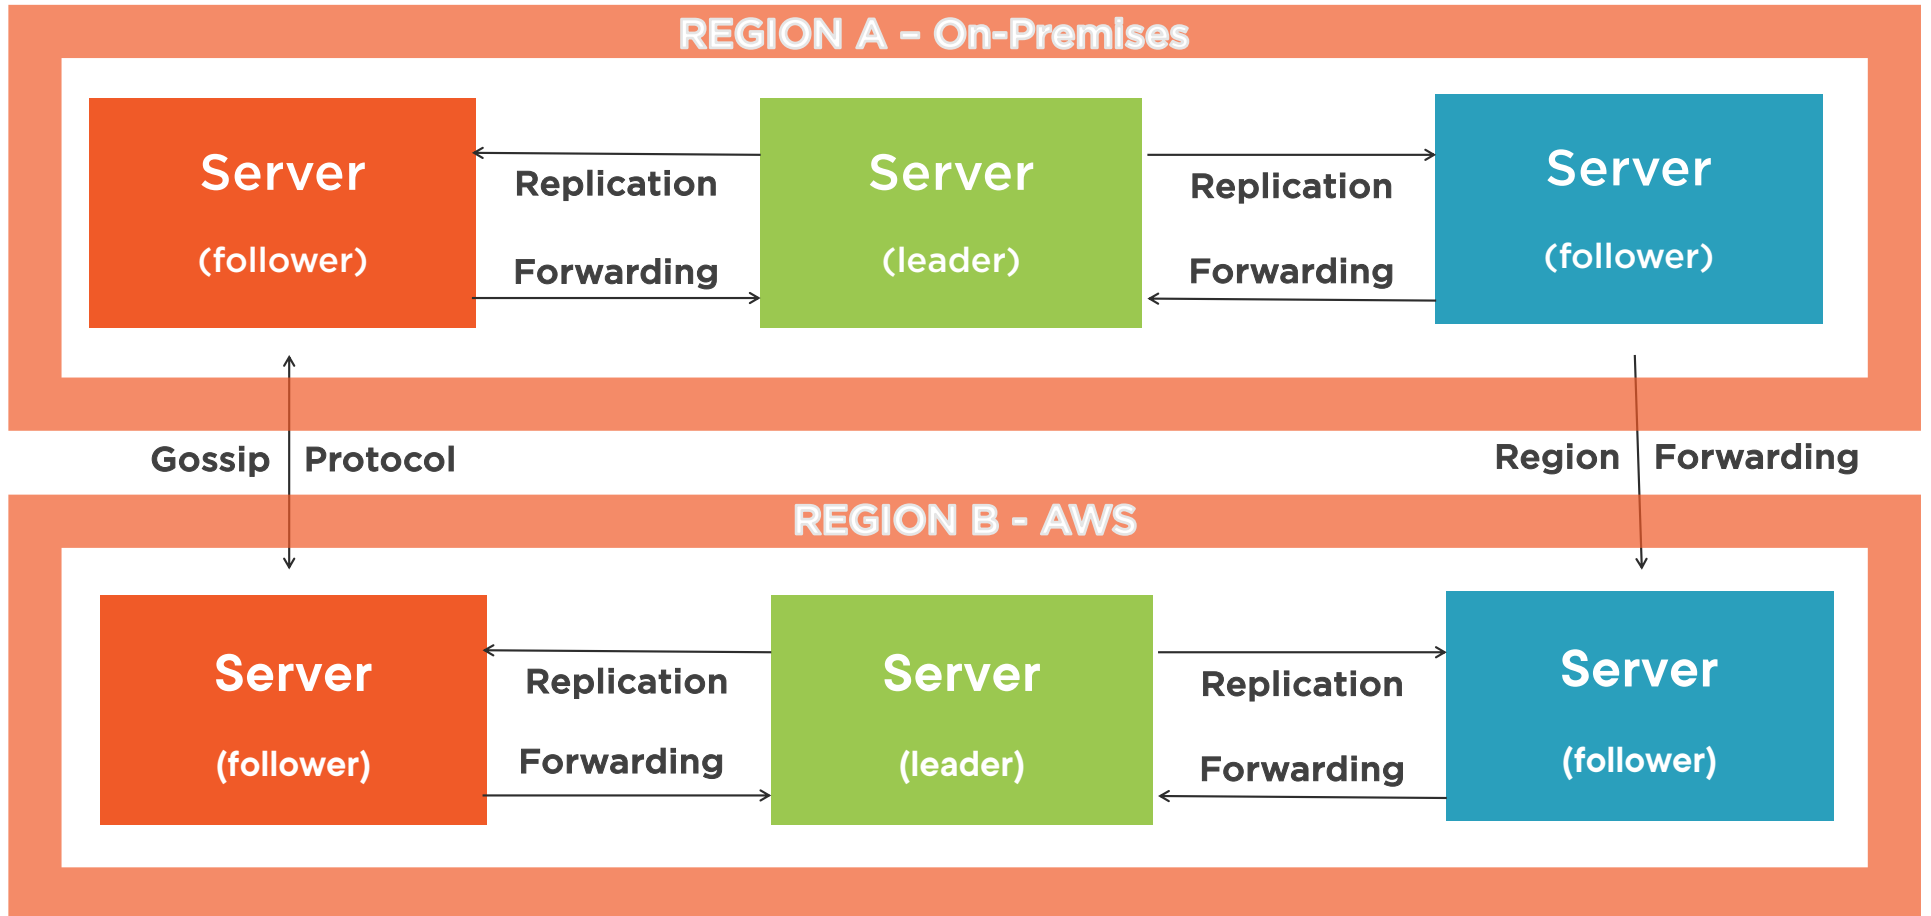
You are running a development environment and need to deploy production using Hybrid infrastructure.

Wired Brain is using bare-metal on-premises and AWS in the cloud with a dedicated persistent VPN connection

# The Wired Brain Hybrid Deployment

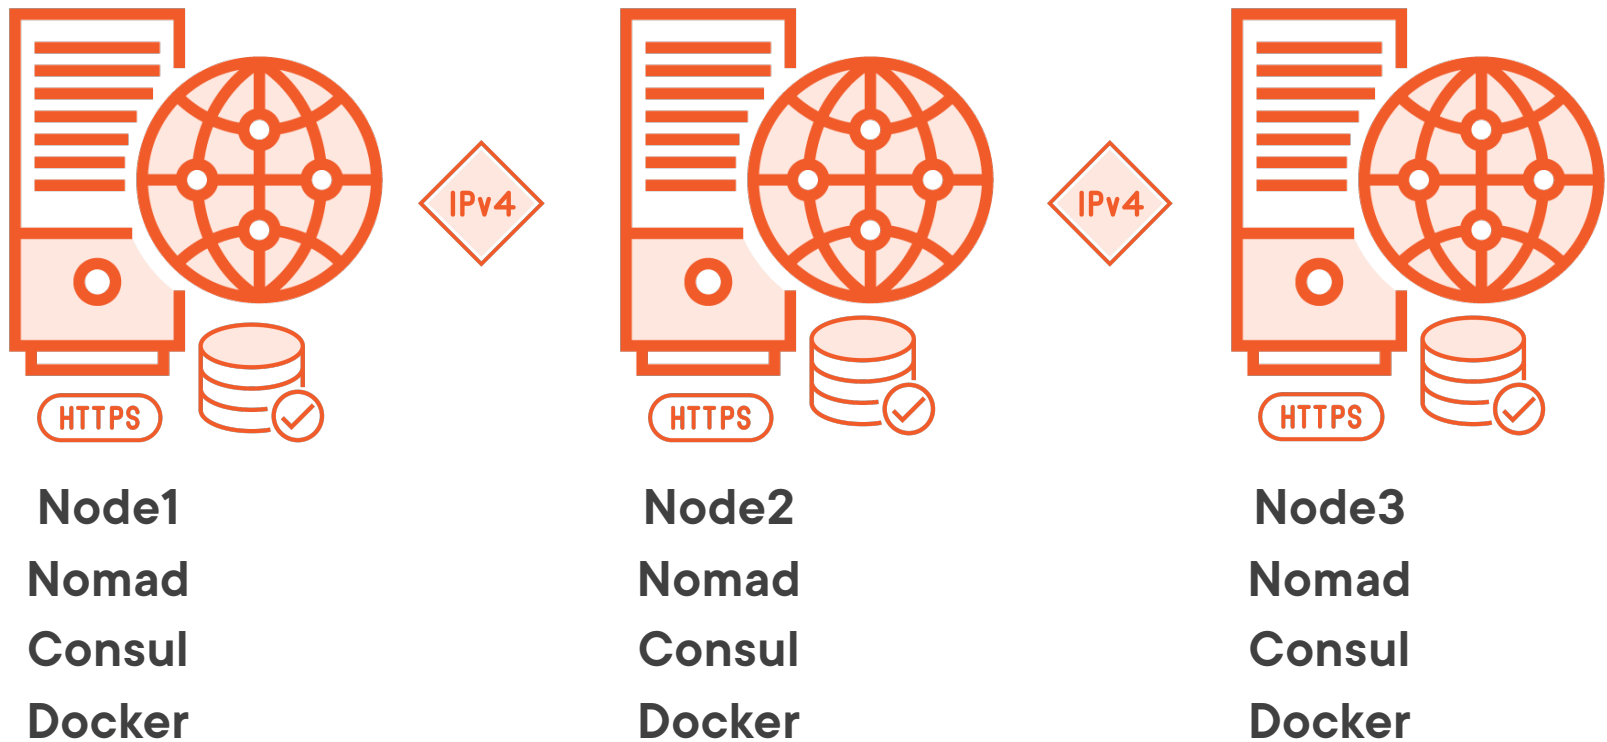
---

# Wired Brain Coffee





## Wired Brain Coffee – Local Data Center



# Nomad “At-Scale”



## The Million Container Challenge


HashiCorp scheduled 1,000,000 Docker containers on 5,000 hosts in under 5 minutes with Nomad, our free and open source cluster scheduler.

<https://www.hashicorp.com/c1m>

# Nomad Deployment Using Infrastructure-as-Code

---

# Demo



- Review the code and configuration for a sample hybrid deployment of Nomad

## Summary



- Reviewed Nomad dependencies
- Explored deployment scenarios
- Chose the Wired Brain Coffee use-case
- Mapped out a hybrid Nomad deployment
- Looked at a code-based deployment of Nomad using Terraform and AWS

Up Next:

Deploying Nomad in a Hybrid Architecture

---